

Introduction to the Course

TIME SERIES ANALYSIS IN PYTHON



Rob Reider


Adjunct Professor, NYU-Courant
Consultant, Quantopian

Quantopian: Business built on python for analyzing backlog of data & quantitative data strategies.
Quantopian is also a community where code is shared and questions are asked/answered.

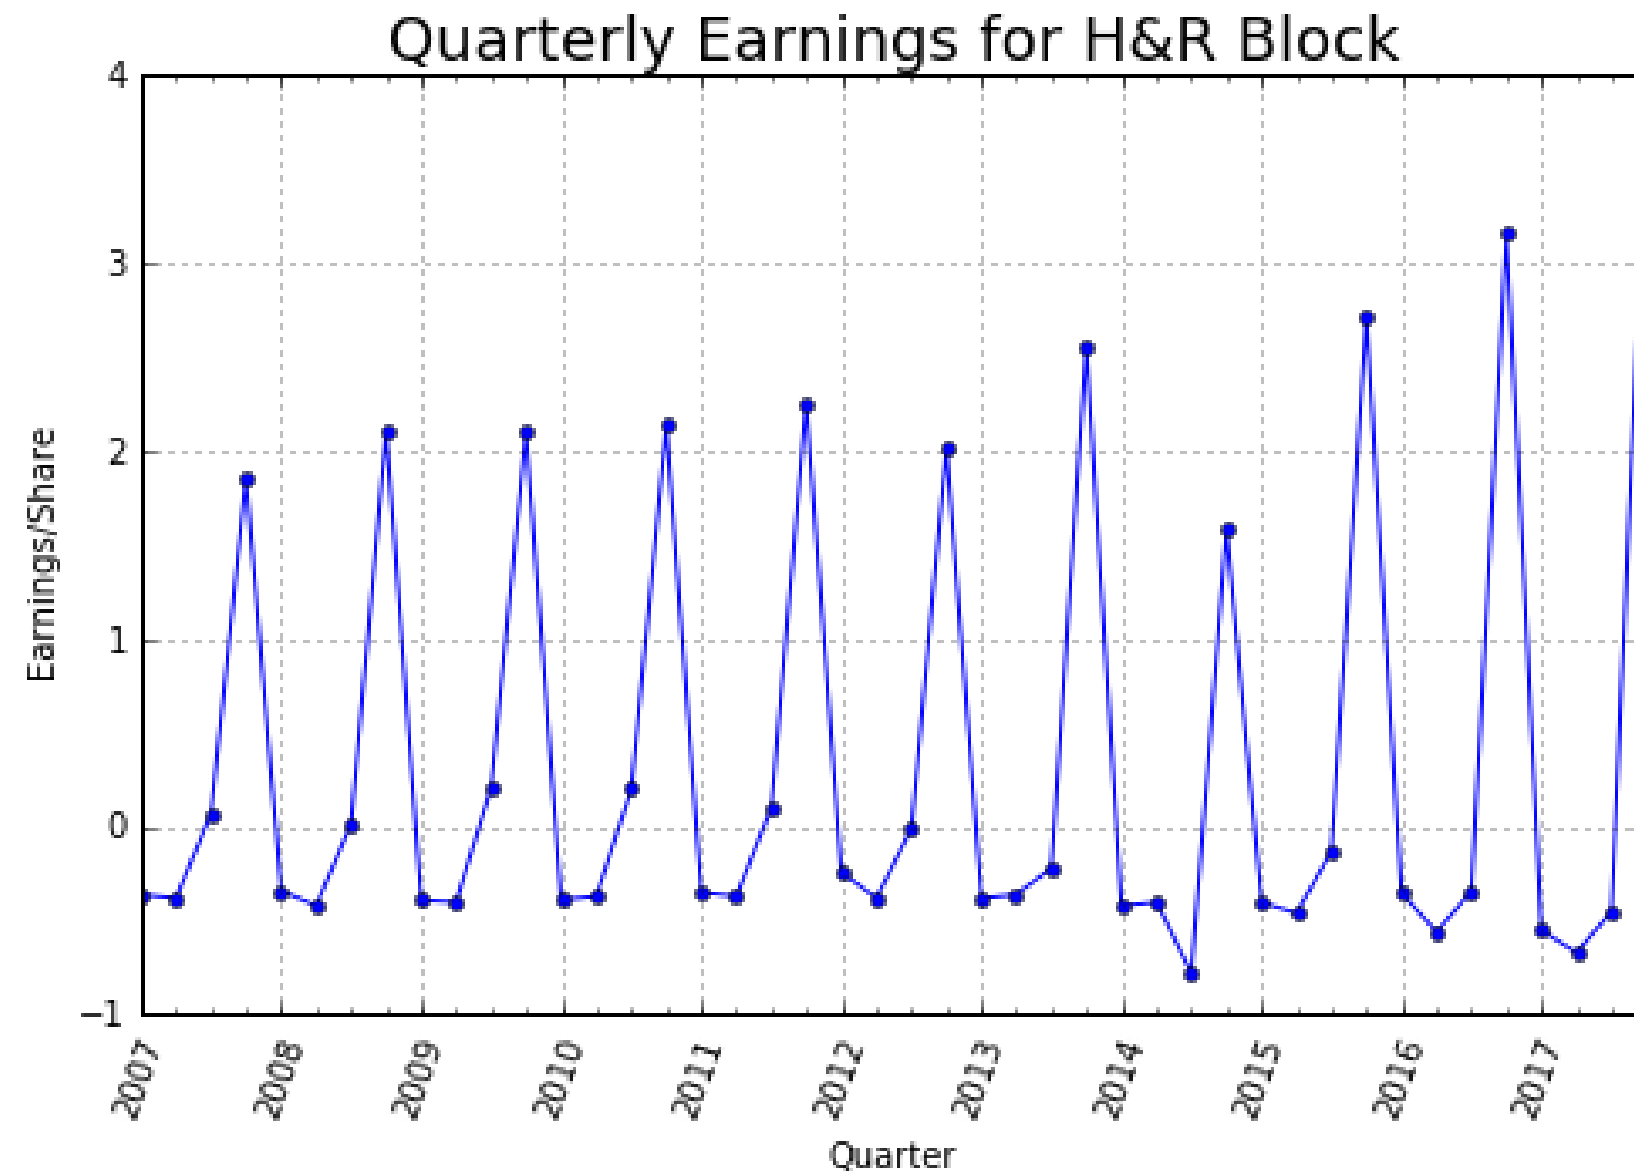
Example of Time Series: Google Trends



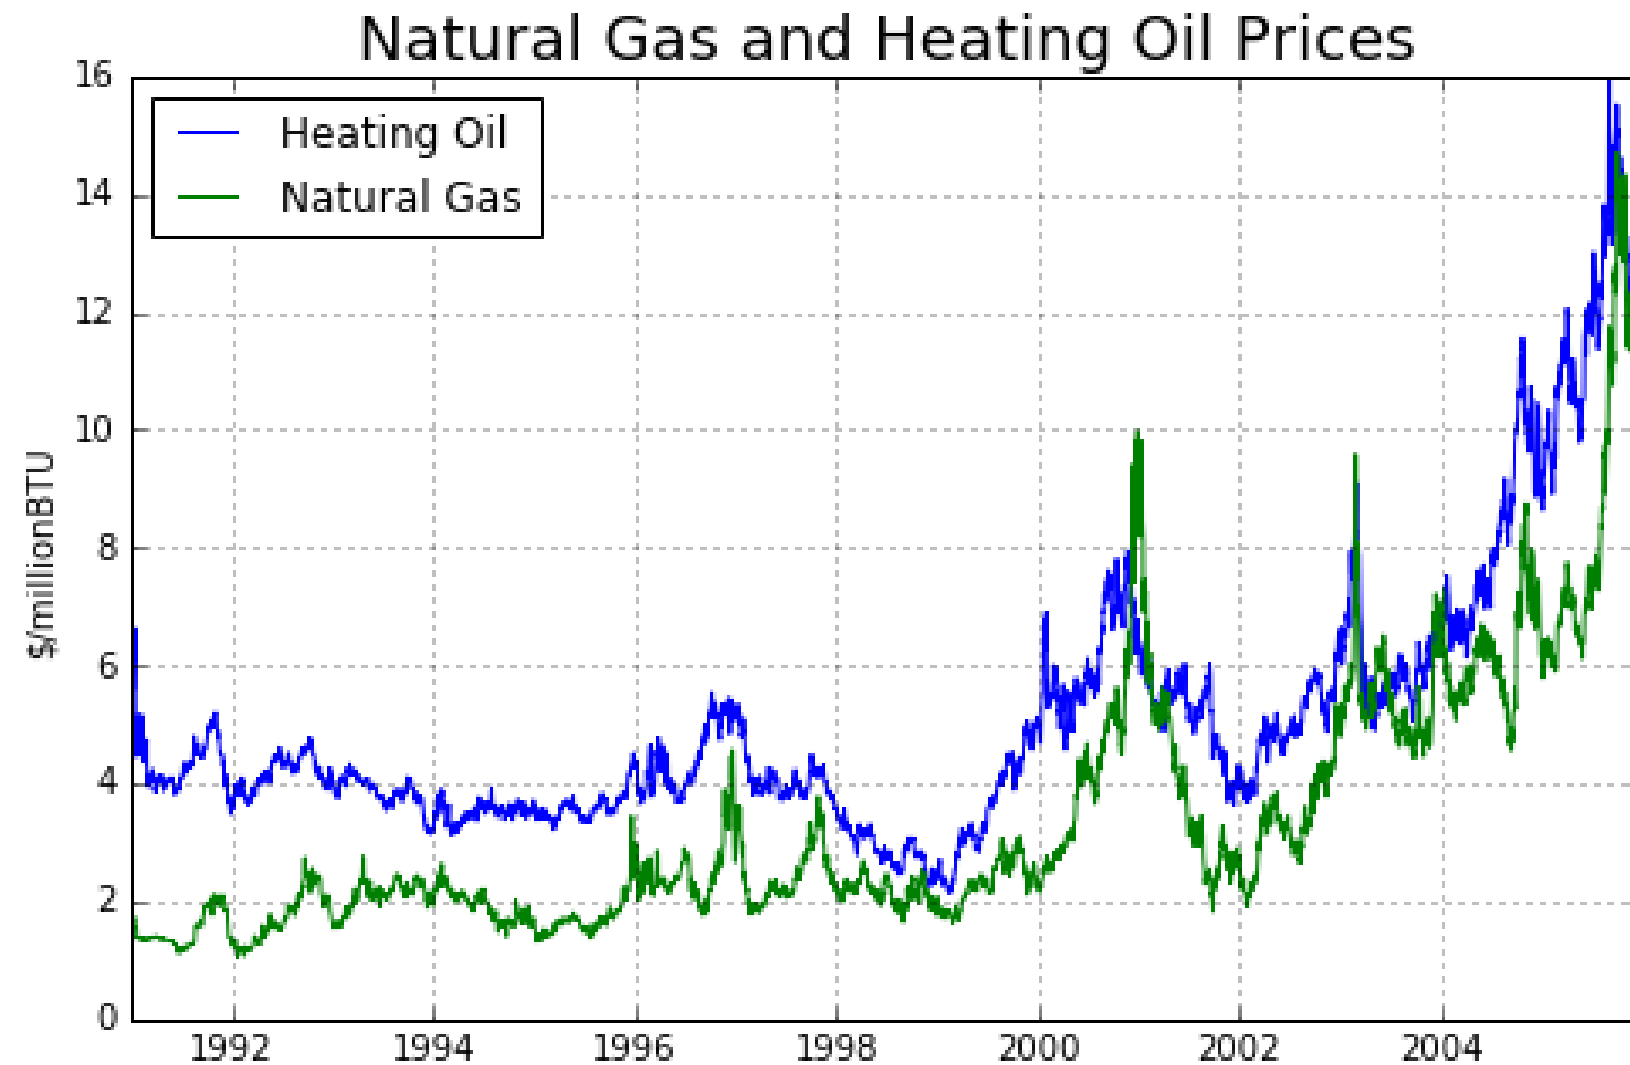
Example of Time Series: Climate Data



Example of Time Series: Quarterly Earnings Data



Example of Multiple Series: Natural Gas and Heating Oil



Goals of Course

- Learn about time series models
- Fit data to a times series model
- Use the models to make forecasts of the future
- Learn how to use the relevant statistical packages in Python
- Provide concrete examples of how these models are used

Many of the models and data will be applicable in the field of finance.

Some Useful Pandas Tools

- Changing an index to datetime

```
df.index = pd.to_datetime(df.index)
```

- Plotting data

```
df.plot()
```

- Slicing data

```
df[ '2012' ]
```

Some Useful Pandas Tools

- Join two DataFrames Say, if one dataframe contains stock prices while another one contains bond prices.

```
df1.join(df2)
```

- Resample data (e.g. from daily to weekly)

```
df = df.resample(rule='W', how='last')
```


More pandas Functions

- Computing percent changes and differences of a time series

```
df[ 'col' ].pct_change() Say, if you wanted to convert prices to returns.  
df[ 'col' ].diff()
```

- pandas correlation method of Series

```
df[ 'ABC' ].corr(df[ 'XYZ' ])
```

- pandas autocorrelation

```
df[ 'ABC' ].autocorr()
```

```
# From previous step
diet.index = pd.to_datetime(diet.index)
```

```
# Slice the dataset to keep only 2012
diet2012 = diet['2012']
```

```
# Plot 2012 data
diet2012.plot(grid = True)
plt.show()
```

```
# Import pandas
import pandas as pd
```

```
# Convert the stock index and bond index into sets
set_stock_dates = set(stocks.index)
set_bond_dates = set(bonds.index)
```

```
# Take the difference between the sets and print
print(set_stock_dates - set_bond_dates)
```

```
# Merge stocks and bonds DataFrames using join()
stocks_and_bonds = stocks.join(bonds, how = 'inner')
```

Let's practice!

TIME SERIES ANALYSIS IN PYTHON

Correlation of Two Time Series

TIME SERIES ANALYSIS IN PYTHON

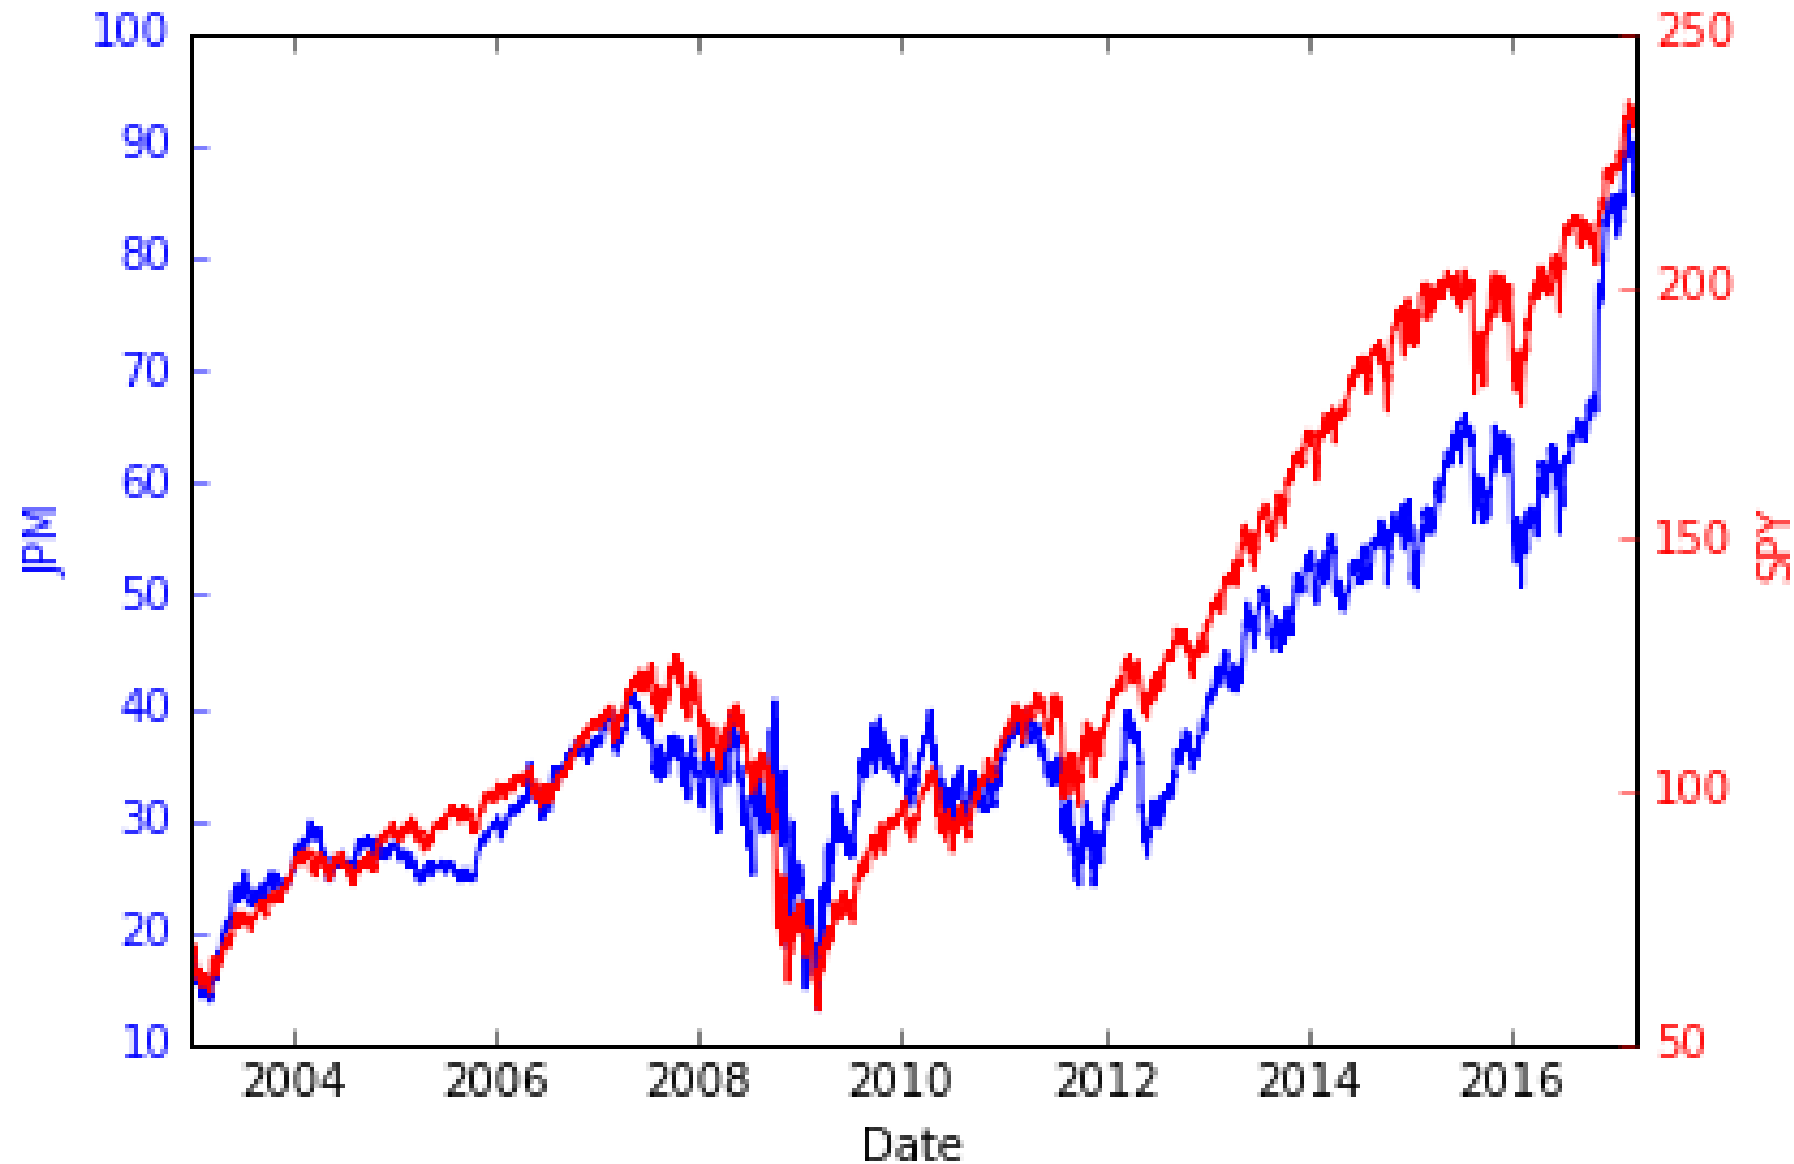


Rob Reider

Adjunct Professor, NYU-Courant
Consultant, Quantopian

Correlation of Two Time Series


- Plot of S&P500 and JPMorgan stock



Here you see that, generally, when the market dips, so does JPM, and when the market (S&P500) rises, so too does JPM.

Correlation of Two Time Series

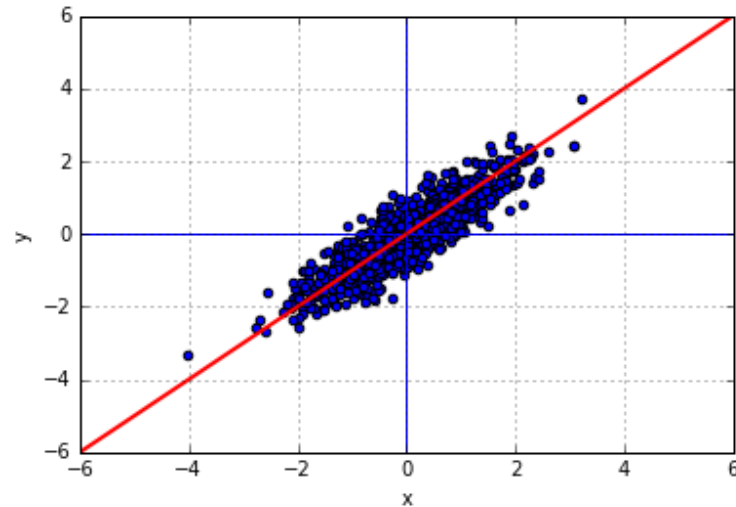
- Scatter plot of S&P500 and JP Morgan returns



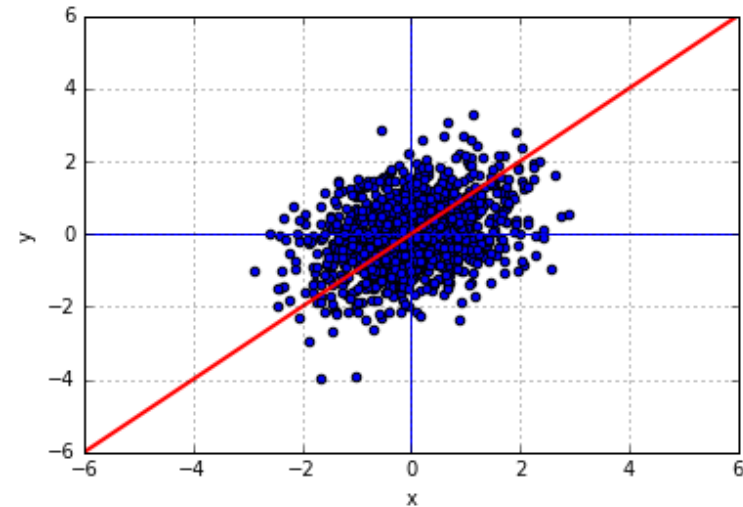
More Scatter Plots

The correlation coefficient is a measure of how much two series vary together.

- Correlation = 0.9



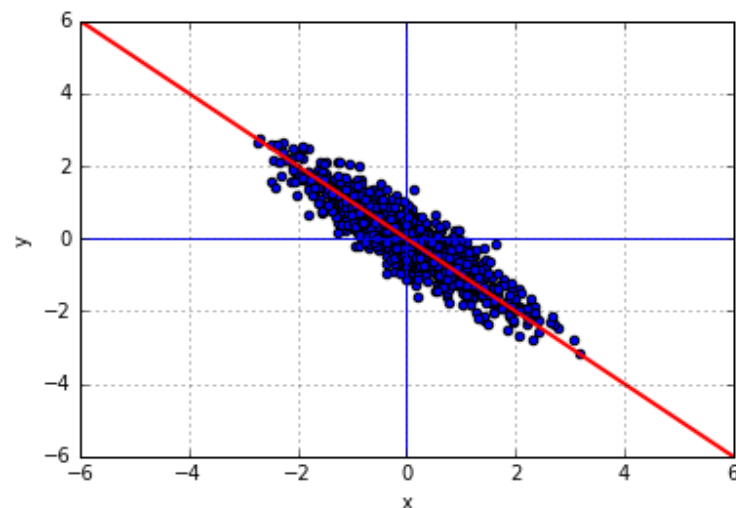
- Correlation = 0.4



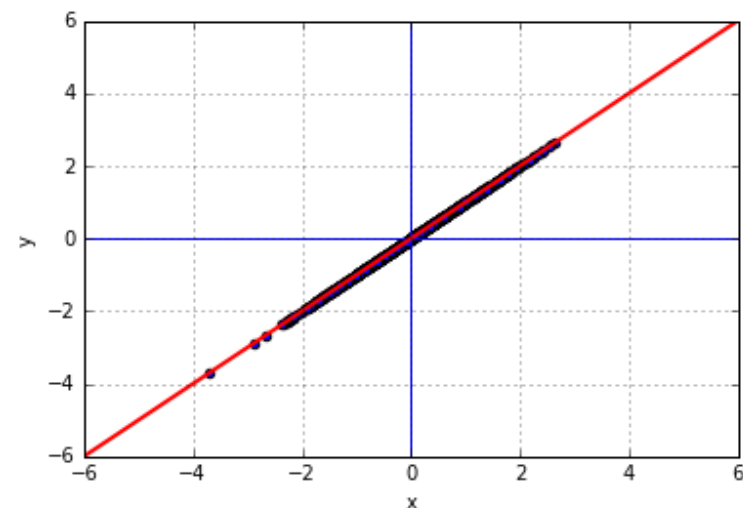
High correlation = the two series strongly vary together.

Low correlation = the two series vary together but with a weak association.

- Correlation = -0.9



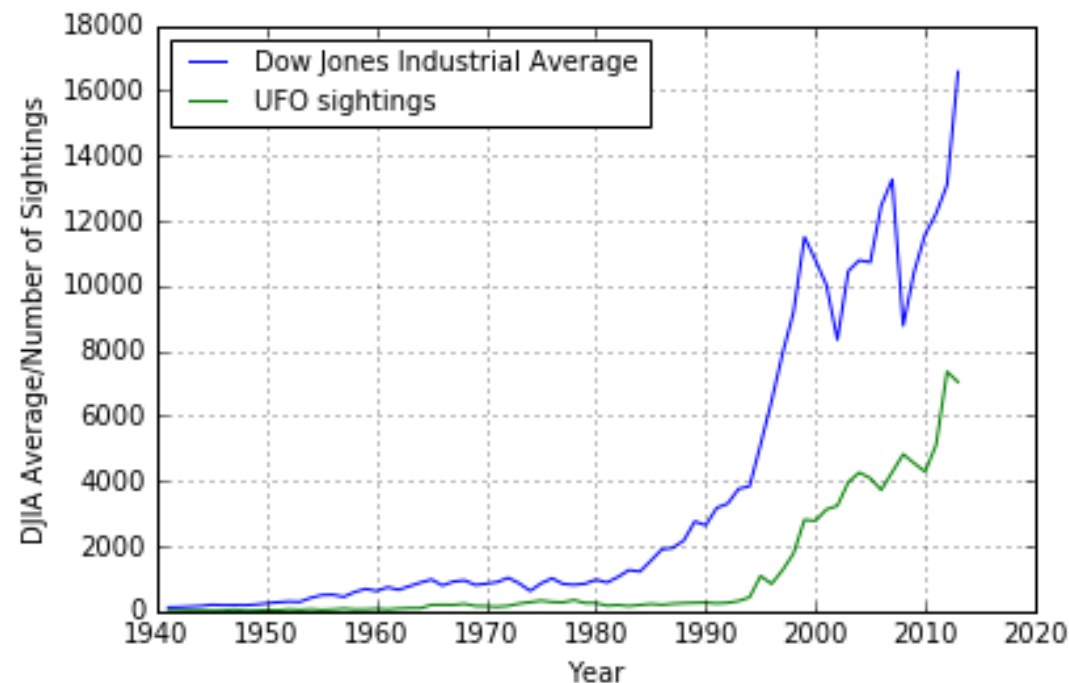
- Correlation = 1.0



High (negative) correlation: the two series vary together, but in opposite directions, but still, with a linear relationship.

Common Mistake: Correlation of Two Trending Series

- Dow Jones Industrial Average and UFO Sightings (www.nuforc.org)



Even if the two series are totally unrelated you can still get a strong correlation.

In the situation where you're looking at the correlation of two stocks, it's important to look at the correlations of their returns rather than their levels.

- Correlation of levels: 0.94

Here, when you look at the correlations of their levels you get the above coefficient, but when you correlate their `pct_change`, you get 0.

Example: Correlation of Large Cap and Small Cap Stocks

- Start with stock prices of SPX (large cap) and R2000 (small cap)
- First step: Compute percentage changes of both series

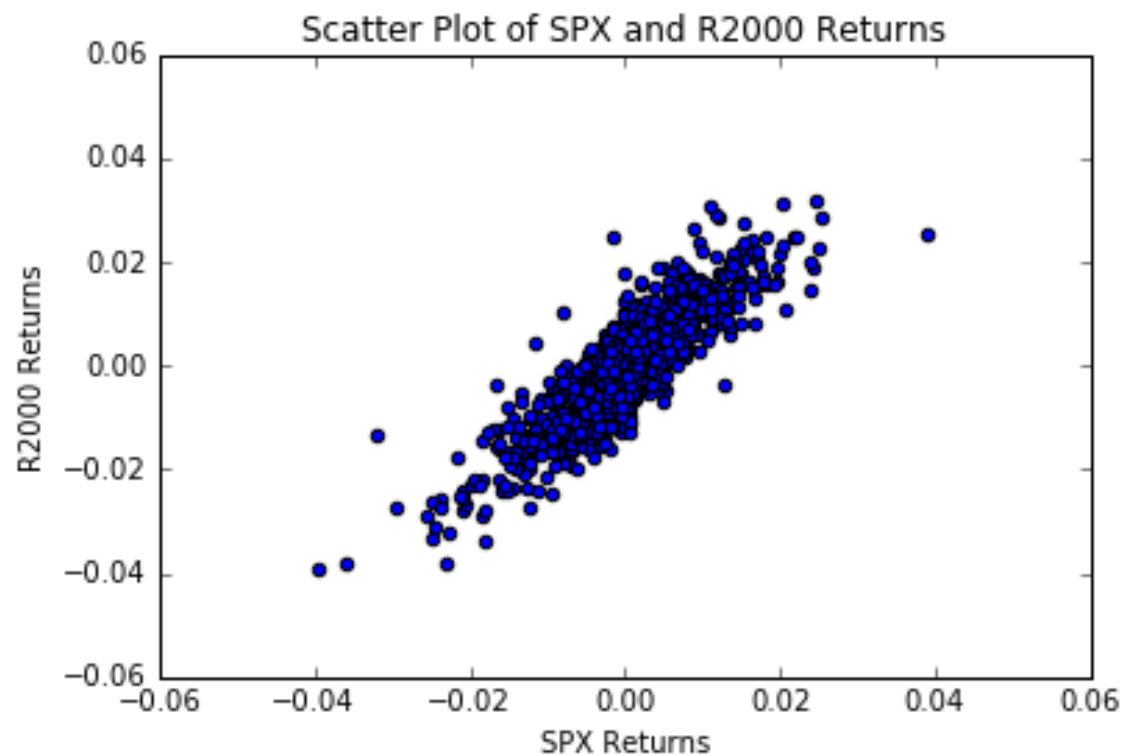
```
df['SPX_Ret'] = df['SPX_Prices'].pct_change()  
df['R2000_Ret'] = df['R2000_Prices'].pct_change()
```

This gives the returns of these series instead of prices.

Example: Correlation of Large Cap and Small Cap Stocks

- Visualize correlation with scatter plot

```
plt.scatter(df['SPX_Return'], df['R2000_Return'])  
plt.show()
```



Example: Correlation of Large Cap and Small Cap Stocks

- Use pandas correlation method for Series

```
correlation = df['SPX_Ret'].corr(df['R2000_Ret'])  
print("Correlation is: ", correlation)
```

```
Correlation is: 0.868
```

```
# Compute percent change using pct_change()
returns = stocks_and_bonds.pct_change()

# Compute correlation using corr()
correlation = returns['SP500'].corr(returns['US10Y'])
print("Correlation of stocks and interest rates: ",
      correlation)

# Make scatter plot
plt.scatter(returns['SP500'],returns['US10Y'])
plt.show()
```

The positive correlation means that when interest rates go down, stock prices go down. For example, during crises like 9/11, investors sold stocks and moved their money to less risky bonds (this is sometimes referred to as a 'flight to quality'). During these periods, stocks drop and interest rates drop as well. Of course, there are times when the opposite relationship holds too.

Let's practice!

TIME SERIES ANALYSIS IN PYTHON

```
# Compute correlation of levels
correlation1 = levels.DJI.corr(levels.UFO)
print("Correlation of levels: ", correlation1)

# Compute correlation of percent changes
changes = levels.pct_change()
correlation2 = changes.DJI.corr(changes.UFO)
print("Correlation of changes: ", correlation2)
```

Notice that the correlation on levels is high but the correlation on changes is close to zero.

Simple Linear Regressions

TIME SERIES ANALYSIS IN PYTHON



Rob Reider

Adjunct Professor, NYU-Courant
Consultant, Quantopian

What is a Regression?


- Simple linear regression: $y_t = \alpha + \beta x_t + \epsilon_t$

A linear regression finds the slope, Beta, and the intercept, Alpha, of a line, that's the best fit between a dependent variable y and an independent variable x.

Xs & Ys can be two time series.

What is a Regression?

- Ordinary Least Squares (OLS)



A linear regression is also known as a Ordinary Least Squares:

It aims to minimize the sum of the squared distances between the data points and the regression line.

Python Packages to Perform Regressions

- In statsmodels:

```
import statsmodels.api as sm
sm.OLS(y, x).fit()
```

Warning: the order of **x** and **y** is not consistent across packages

- In numpy:

```
np.polyfit(x, y, deg=1)
```

← — setting 'deg' parameter to 1 is the command for a linear regression.

- In pandas:

```
pd.ols(y, x)
```

- In scipy:

```
from scipy import stats
stats.linregress(x, y)
```

Beware, that the order of x & y is not consistent across packages.

Example: Regression of Small Cap Returns on Large Cap

- Import the statsmodels module

```
import statsmodels.api as sm
```

- As before, compute percentage changes in both series

```
df['SPX_Ret'] = df['SPX_Prices'].pct_change()  
df['R2000_Ret'] = df['R2000_Prices'].pct_change()
```

- Add a constant to the DataFrame for the regression intercept

```
df = sm.add_constant(df)
```

If you don't have a constant column then the packages assumed you don't want to run the regression with an intercept.

By adding a column of 1s, statsmodels will compute the regression coefficient of that column as well, which can be interpreted as the intercept of the line.

Regression Example (continued)

- Notice that the first row of returns is NaN

	SPX_Price	R2000_Price	SPX_Ret	R2000_Ret
Date				
2012-11-01	1427.589966	827.849976	NaN	NaN
2012-11-02	1414.199951	814.369995	-0.009379	-0.016283

Because the return is calculated from a previous price, the first line has no previous price and therefore can't produce a `pct_change()` value.

- Delete the row of NaN

```
df = df.dropna()
```

- Run the regression

```
results = sm.OLS(df['R2000_Ret'], df[['const', 'SPX_Ret']]).fit()  
print(results.summary())
```

The first argument is your dependent variable `y`, the next argument is your independent variable or variables for `x`.

Regression Example (continued)

- Regression output

```
=====
                        OLS Regression Results
=====
Dep. Variable:          R2000_Ret      R-squared:                0.753
Model:                  OLS           Adj. R-squared:           0.753
Method:                 Least Squares   F-statistic:              3829.
Date:                   Fri, 26 Jan 2018 Prob (F-statistic):       0.00
Time:                   13:29:55        Log-Likelihood:          4882.4
No. Observations:       1257           AIC:                    -9761.
Df Residuals:           1255           BIC:                    -9751.
Df Model:                1
Covariance Type:        nonrobust
=====

```

	coef	std err	t	P> t	[95.0% Conf. Int.]
const	-4.964e-05	0.000	-0.353	0.724	-0.000 0.000
SPX_Ret	1.1412	0.018	61.877	0.000	1.105 1.177

```
=====
Omnibus:                 61.950      Durbin-Watson:           1.991
Prob(Omnibus):            0.000      Jarque-Bera (JB):        148.100
Skew:                     0.266      Prob(JB):                6.93e-33
Kurtosis:                 4.595      Cond. No.:               131.
=====
```

const is the intercept or Alpha
SPX_Ret is the Slope or Beta

- Intercept in `results.params[0]`
- Slope in `results.params[1]`

Regression Example (continued)

- Regression output

```

=====
                        OLS Regression Results
=====
Dep. Variable:          R2000_Ret      R-squared:                0.753
Model:                  OLS            Adj. R-squared:           0.753
Method:                 Least Squares   F-statistic:              3829.
Date:                   Fri, 26 Jan 2018 Prob (F-statistic):       0.00
Time:                   13:29:55        Log-Likelihood:          4882.4
No. Observations:       1257           AIC:                    -9761.
Df Residuals:           1255           BIC:                    -9751.
Df Model:                1
Covariance Type:        nonrobust
=====

```

	coef	std err	t	P> t	[95.0% Conf. Int.]
const	-4.964e-05	0.000	-0.353	0.724	-0.000 0.000
SPX_Ret	1.1412	0.018	61.877	0.000	1.105 1.177

```

=====
Omnibus:                 61.950   Durbin-Watson:              1.991
Prob(Omnibus):            0.000   Jarque-Bera (JB):          148.100
Skew:                     0.266   Prob(JB):                  6.93e-33
Kurtosis:                  4.595   Cond. No.                   131.
=====

```

Relationship Between R-Squared and Correlation

- $[\text{corr}(x, y)]^2 = R^2$ (or R-squared)
- $\text{sign}(\text{corr}) = \text{sign}(\text{regression slope})$
- In last example:
 - R-Squared = 0.753
 - Slope is positive
 - $\text{correlation} = +\sqrt{0.753} = 0.868$

In the same way that a correlation measures how closely the data are clustered along a line, so too does R-squared.

R-squared measures how well the linear regression line fits the data.

There is a relationship between correlation and r-squared: The magnitude of the correlation is the square root of the r-squared. The sign of the correlation is the sign of the slope or Beta

If the regression line is positive then the correlation is positive and vice-versa.

```
# Import the statsmodels module
import statsmodels.api as sm

# Compute correlation of x and y
correlation = x.corr(y)
print("The correlation between x and y is %4.2f" %(correlation))

# Convert the Series x to a DataFrame and name the column x
dfx = pd.DataFrame(x, columns=['x'])

# Add a constant to the DataFrame dfx
dfx1 = sm.add_constant(dfx)

# Regress y on dfx1
result = sm.OLS(y, dfx1).fit()

# Print out the results and look at the relationship between R-squared and the correlation above
print(result.summary())
```


Notice that the two different methods of computing correlation give the same result. The correlation is about -0.9 and the R-squared is about 0.81

Let's practice!

TIME SERIES ANALYSIS IN PYTHON

Match Correlation with Regression Output

Here are four scatter plots, each showing a linear regression line and an R-squared:



Which correlation is correct?

✓ Answer the question

50 XP

Possible Answers

- ☐ Fig 1: correlation = -0.6
- ☐ Fig 2: correlation = -0.9
- ☒ Fig 3: correlation = -0.9
- ☐ Fig 4: correlation = -0.32

💡 Take Hint (-15xp)

press 1

press 2

press 3

press 4

Submit Answer

Autocorrelation

TIME SERIES ANALYSIS IN PYTHON



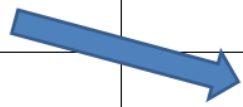
Rob Reider

Adjunct Professor, NYU-Courant
Consultant, Quantopian

What is Autocorrelation?

- Correlation of a time series with a lagged copy of itself


Series	Lagged Series
5	
10	5
15	10
20	15
25	20
⋮	⋮



- Lag-one autocorrelation
- Also called **serial correlation**


Interpretation of Autocorrelation

- **Mean Reversion - Negative autocorrelation** <— In financial time series, this is the case, also expressed as Mean Reverting.



Interpretation of Autocorrelation

- Momentum, or Trend Following - Positive autocorrelation



Traders Use Autocorrelation to Make Money

- Individual stocks
 - Historically have negative autocorrelation
 - Measured over short horizons (days)
 - Trading strategy: Buy losers and sell winners
- Commodities and currencies
 - Historically have positive autocorrelation
 - Measured over longer horizons (months)
 - Trading strategy: Buy winners and sell losers

Example of Positive Autocorrelation: Exchange Rates

- Use daily ¥/\$ exchange rates in DataFrame `df` from **FRED** <--- Federal Reserve Economic Data
- Convert index to datetime

```
# Convert index to datetime
df.index = pd.to_datetime(df.index)
# Downsample from daily to monthly data
df = df.resample(rule='M', how='last')
# Compute returns from prices
df['Return'] = df['Price'].pct_change()
# Compute autocorrelation
autocorrelation = df['Return'].autocorr()
print("The autocorrelation is: ", autocorrelation)
```

```
The autocorrelation is: 0.0567
```

Because the autocorrelation is positive, the series is said to have 'momentum'

```
# Convert the daily data to weekly data
MSFT = MSFT.resample(rule = 'W').last()

# Compute the percentage change of prices
returns = MSFT.pct_change()

# Compute and print the autocorrelation of returns
autocorrelation = returns['Adj Close'].autocorr()
print("The autocorrelation of weekly returns is %4.2f" %
(autocorrelation))

Notice how the autocorrelation of returns for MSFT is negative, so the
stock is 'mean reverting'
```

Let's practice!

TIME SERIES ANALYSIS IN PYTHON

```
# Compute the daily change in interest rates
daily_diff = daily_rates.diff()

# Compute and print the autocorrelation of daily changes
autocorrelation_daily = daily_diff['US10Y'].autocorr()
print("The autocorrelation of daily interest rate changes is %4.2f" %(autocorrelation_daily))

# Convert the daily data to annual data
yearly_rates = daily_rates.resample(rule = 'A').last()

# Repeat above for annual data
yearly_diff = yearly_rates.diff()
autocorrelation_yearly = yearly_diff['US10Y'].autocorr()
print("The autocorrelation of annual interest rate changes is %4.2f" %(autocorrelation_yearly))

Notice how the daily autocorrelation is small but the annual autocorrelation is large and negative
```